

Machine Learning

Lecture 11: Embeddings

Iurii Efimov



girafe
ai

Outline



1. NLP introduction
2. Text preprocessing
3. Feature extraction:
 - a. Bag-of-Words
 - b. Bag-of-Ngrammes
 - c. TF-IDF
4. Word embeddings
5. Word2vec

Natural Language Processing

girafe
ai


01

Natural language processing

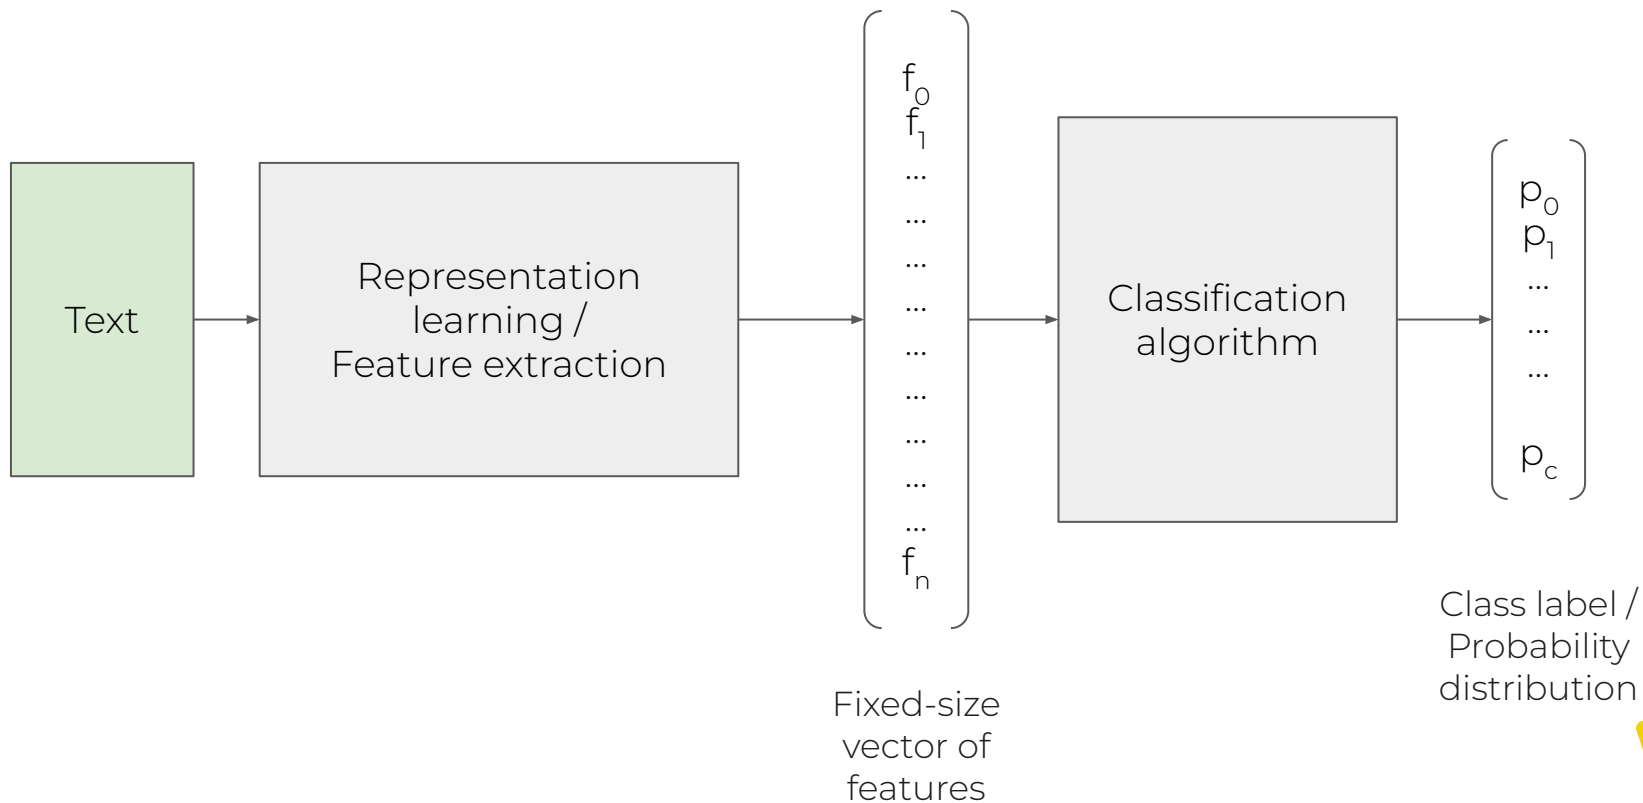


- Text classification tasks:
 - Sentiment analysis
 - Spam filtering
 - Fake news detection
 - Topic prediction
 - #hashtag prediction

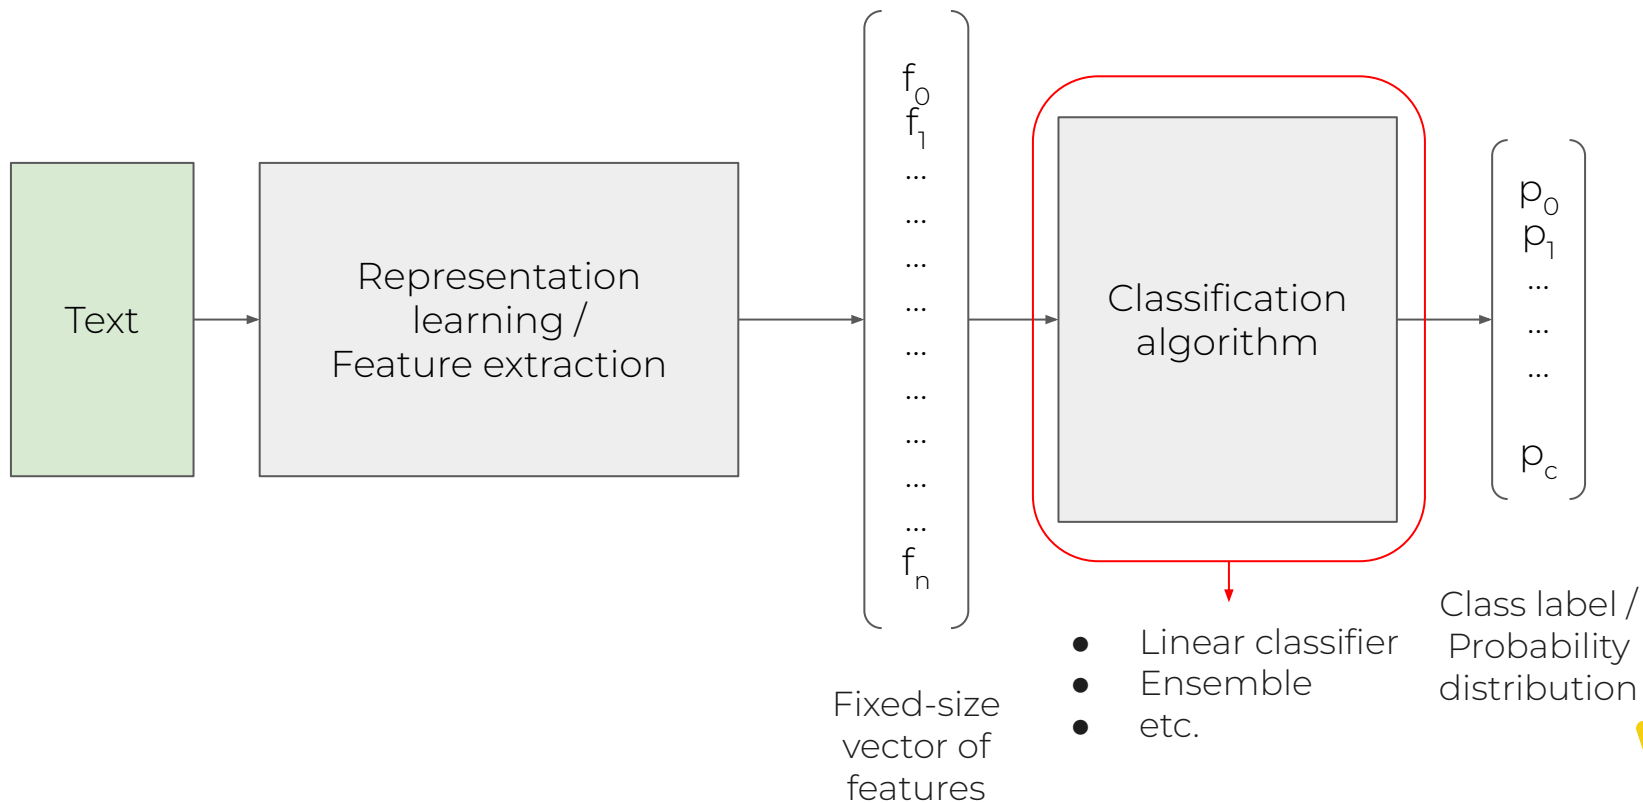
Example: sentiment analysis



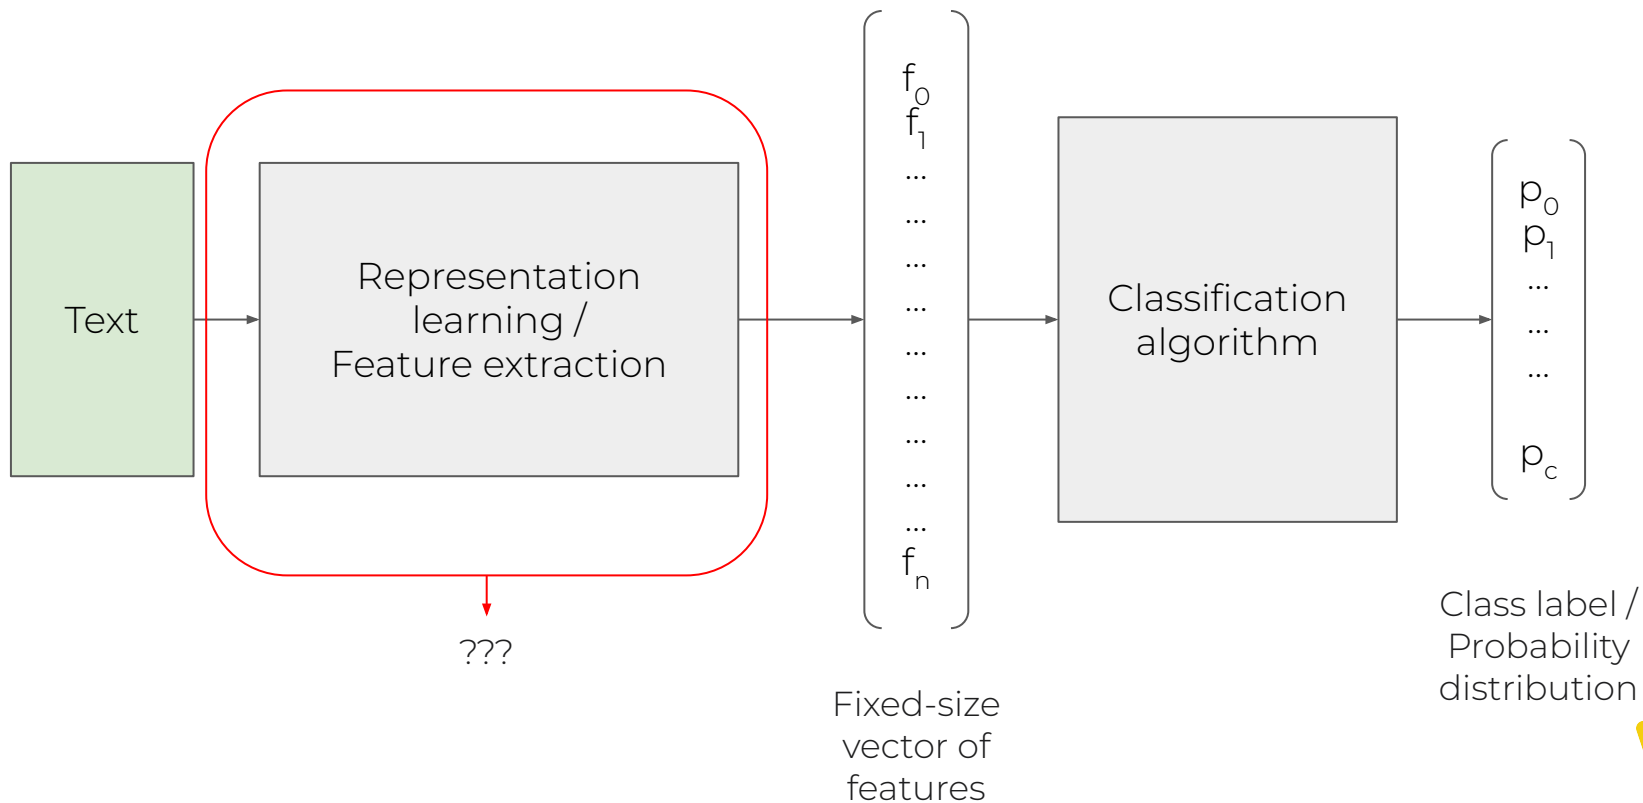
Text classification in general



Text classification in general



Text classification in general



Text preprocessing


girafe
ai

02




Tokenization

- Split the input into tokens
- Tokens:
 - words
 - symbol
 - morpheme
 - ...



Token normalization






Token normalization

Stemming:

- removing and replacing suffixes
- get to the root of the word (stem)






Token normalization

Lemmatization:

- Get base or dictionary form of a word (lemma)



Stemming: Porter vs Lancaster



Porter stemmer

- Published in 1979
- Base starting option

Lancaster stemmer

- Published in 1990
- The most aggressive
- Easy adding of your own rules

Snowball stemmer (Porter 2)

- Based on Porter
- More aggressive
- Most popular option now



Stemming: Example

Porter's stemmer:


- **Heuristics, applied one-by-one:**
 - SSES - SS (dresses - dress)
 - IES - I (ponies - poni)
 - S - <empty> (dogs - dog)
- **What's wrong?**
 - Overstemming and understemming



Overstemming

- University
 - Universal
 - Universities
 - Universe
- } Univers

Understemming

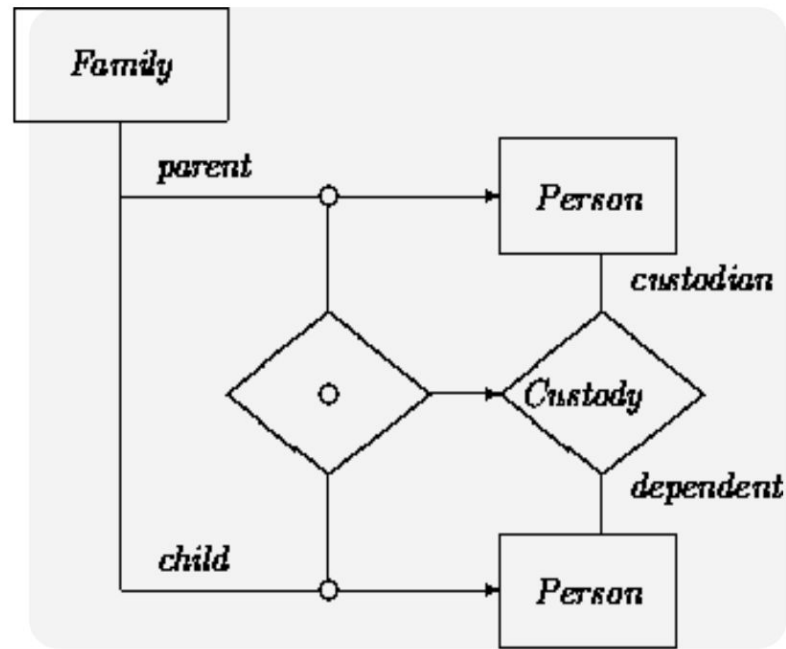
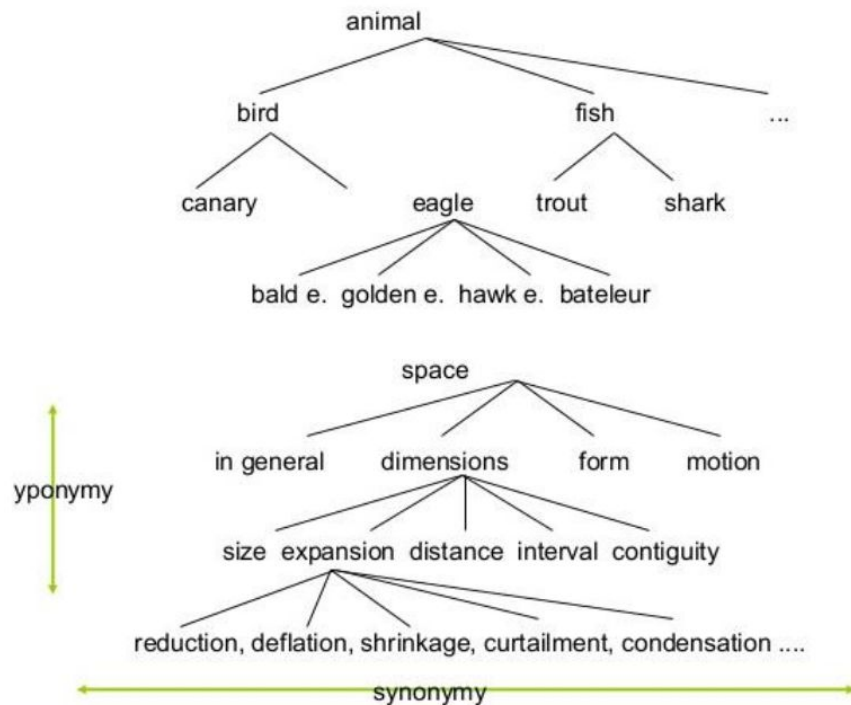


Lemmatization



- Lemmatizer from NLTK:
 - Tries to resolve word to its dictionary form
 - Based on **WordNet** database
 - For the best results feed part-of-speech tagger

What is WordNet?





Handful tools for preprocessing

- [NLTK](#)
 - nltk.stem.SnowballStemmer
 - nltk.stem.PorterStemmer
 - nltk.stem.WordNetLemmatizer
 - nltk.corpus.stopwords
- [BeautifulSoup \(for parsing HTML\)](#)
- Regular Expressions (import re)
- [Pymorphy2](#)

What's left



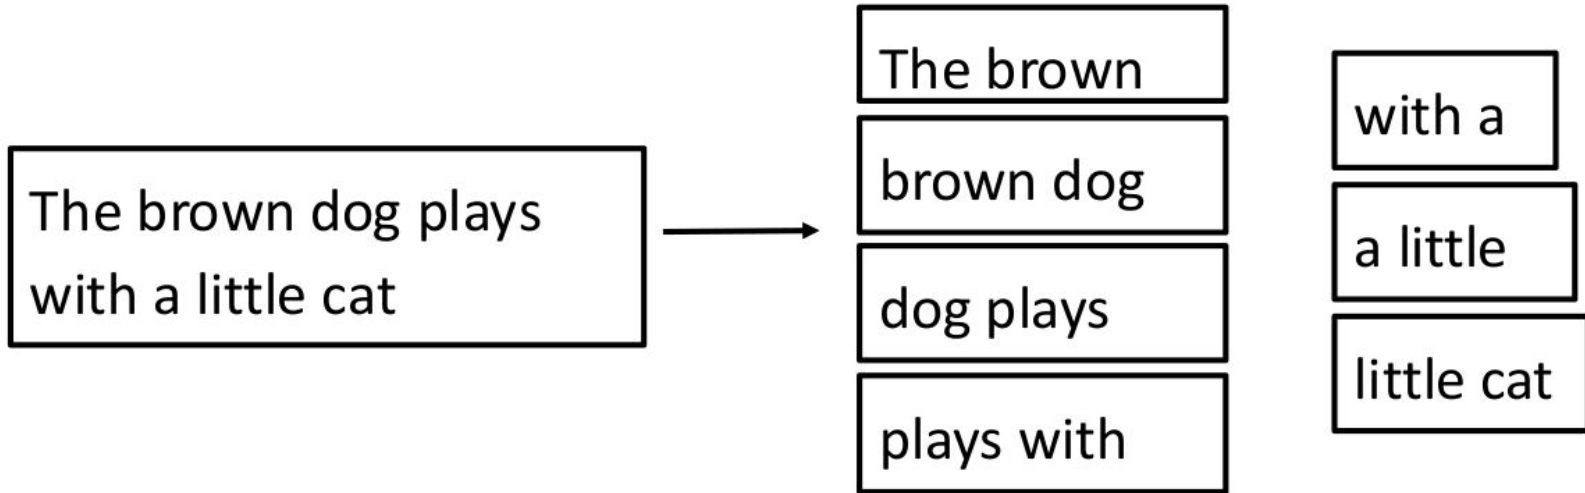
- Capital Letters
- Punctuation
- Contractions (e.g, etc.)
- Numbers (dates, ids, page numbers)
- Stop-words (“the”, “is”, etc.)
- Tags

Bag-of-words

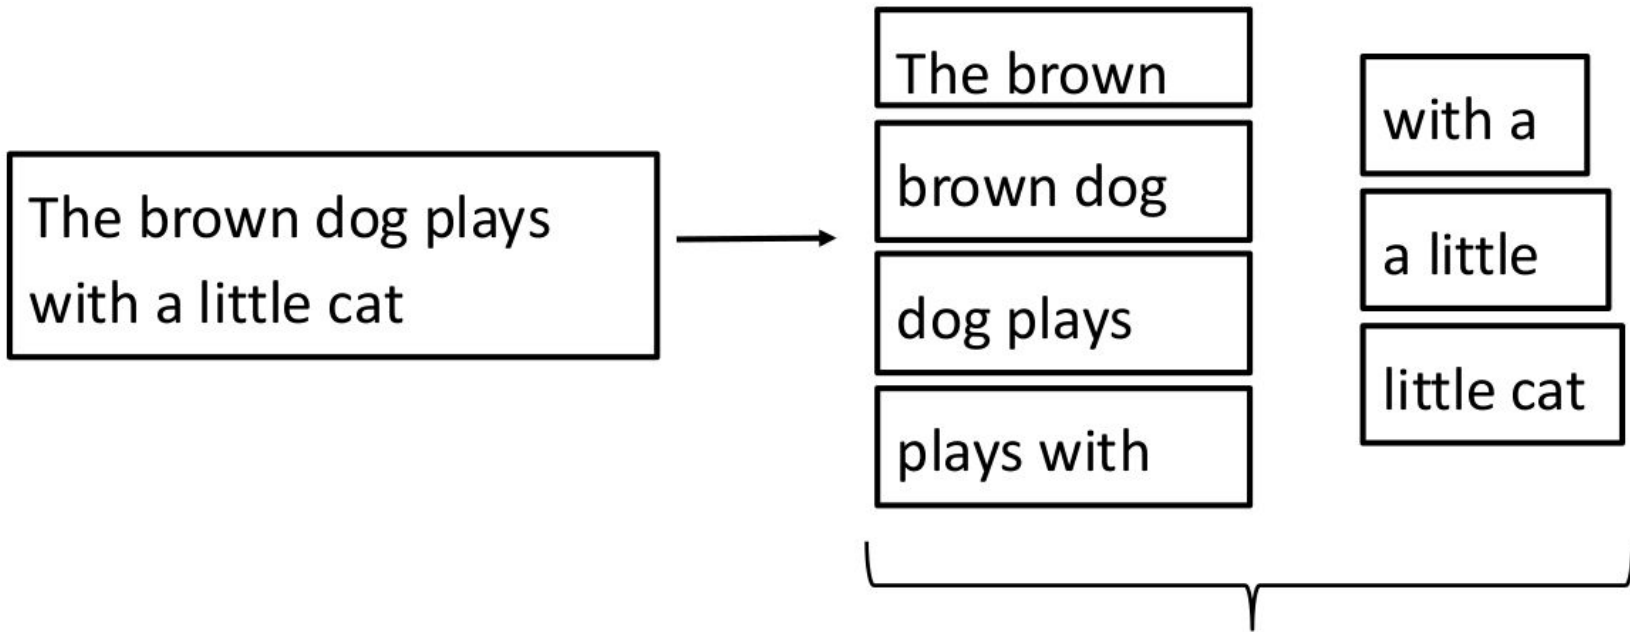


How to improve BOW?

- Use n-gramms instead of words!

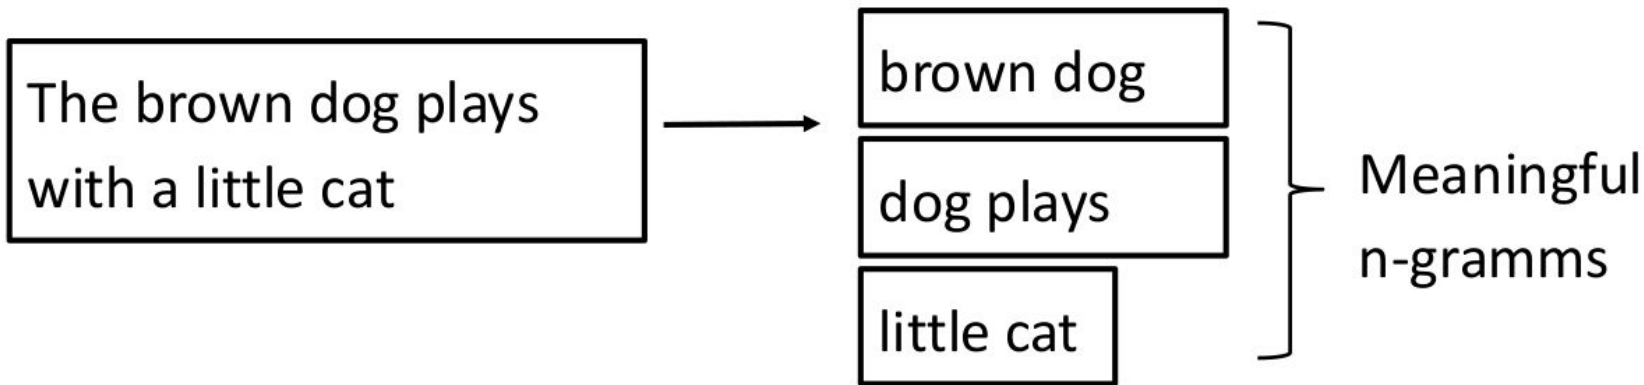


Bag-of-words



Do we need all this bigramms?

Bag-of-words



Meaningful n-gramms are often called **collocations**

How to detect meaningful n-gramms?



Collocations: first step

- **Delete:**

- High-frequency n-gramms
 - Articles, prepositions
 - Auxiliary verbs (to be, to have, etc.)
 - General vocabulary
- Low-frequency n-gramms
 - Typos
 - Combinations that occur 1-2 times in a text

Feature extraction

girafe
ai

03



Bag-of-words

the dog is on the table

0	0	1	1	0	1	1	1
are	cat	dog	is	now	on	table	the

Problems

- No information about words order
- Word vectors are huge and sparse
- Word vectors are not normalized
- Same words can take different forms
 -



TF-IDF

- **Term Frequency (tf):** gives us the frequency of the word in each document in the corpus.

$$\text{tf}(t, d) = f_{t, d}$$

- **Inverse Document Frequency (idf):** used to calculate the weight of rare words across all documents in the corpus. The words that occur rarely in the corpus have a high IDF score.

$$\text{idf}(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

N : total number of documents in the corpus $N = |D|$

$|\{d \in D : t \in d\}|$: number of documents where the term t



TF-IDF: Example

- **Sentence A:** The car is driven on the road.
- **Sentence B:** The truck is driven on the highway

(each sentence is a separate document)

TF-IDF: Example



Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7			
Car	1/7	0			
Truck	0	1/7			
Is	1/7	1/7			
Driven	1/7	1/7			
On	1/7	1/7			
The	1/7	1/7			
Road	1/7	0			
Highway	0	1/7			

TF-IDF: Example



Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$		
Car	1/7	0	$\log(2/1)=0.3$		
Truck	0	1/7	$\log(2/1)=0.3$		
Is	1/7	1/7	$\log(2/2)=0$		
Driven	1/7	1/7	$\log(2/2)=0$		
On	1/7	1/7	$\log(2/2)=0$		
The	1/7	1/7	$\log(2/2)=0$		
Road	1/7	0	$\log(2/1)=0.3$		
Highway	0	1/7	$\log(2/1)=0.3$		

TF-IDF: Example



Word	TF		IDF	TF * IDF	
	A	B		A	B
The	1/7	1/7	$\log(2/2)=0$	0	0
Car	1/7	0	$\log(2/1)=0.3$	0.043	0
Truck	0	1/7	$\log(2/1)=0.3$	0	0.043
Is	1/7	1/7	$\log(2/2)=0$	0	0
Driven	1/7	1/7	$\log(2/2)=0$	0	0
On	1/7	1/7	$\log(2/2)=0$	0	0
The	1/7	1/7	$\log(2/2)=0$	0	0
Road	1/7	0	$\log(2/1)=0.3$	0.043	0
Highway	0	1/7	$\log(2/1)=0.3$	0	0.043

TF-IDF: Example



```
from sklearn.feature_extraction.text import TfidfVectorizer
```



Word Embeddings

girafe
ai

04

One-hot vectors



Problems:

- Huge vectors
- VERY sparse
- No semantics or word similarity information included

Rome Paris word V

Rome = [1, 0, 0, 0, 0, 0, ..., 0]

Paris = [0, 1, 0, 0, 0, 0, ..., 0]

Italy = [0, 0, 1, 0, 0, 0, ..., 0]

France = [0, 0, 0, 1, 0, 0, ..., 0]

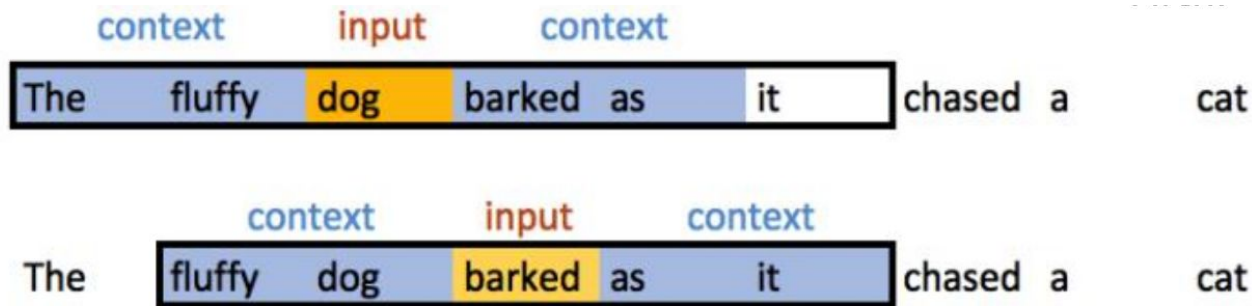


Distributional semantics

Does vector similarity imply semantic similarity?

“You shall know a word by the company it keeps”


Firth





Matrix factorization

- Input: PMI, word cooccurrences, etc.
- Method: dimensionality reduction (SVD)
- Output: word similarities





Collocations: first step

- **Delete:**

- High-frequency n-gramms
 - Articles, prepositions
 - Auxiliary verbs (to be, to have, etc.)
 - General vocabulary
- Low-frequency n-gramms
 - Typos
 - Combinations that occur 1-2 times in a text

Collocations: context is all you need



- Co-Occurrence counters in a window of fixed size
 - states for the number of times we've seen word u and word v together in the window
- Better solution: Pointwise Mutual Information (PMI)

$$PMI = \log \frac{p(u, v)}{p(u)p(v)} = \log \frac{n_{uv}n}{n_u n_v}$$

- Much better solution: Positive PMI (pPMI)

$$pPMI = \max(0, PMI)$$

Learning word vectors

girafe
ai

05

Embeddings: intuition




What is king - man + woman?



Embeddings: intuition




So $\text{king} - \text{man} + \text{woman} = \text{queen!}$




Word2vec




Word2vec (Mikolov et al. 2013) - a framework for learning word embeddings



Male-Female



Verb tense



Country-Capital

Word2vec

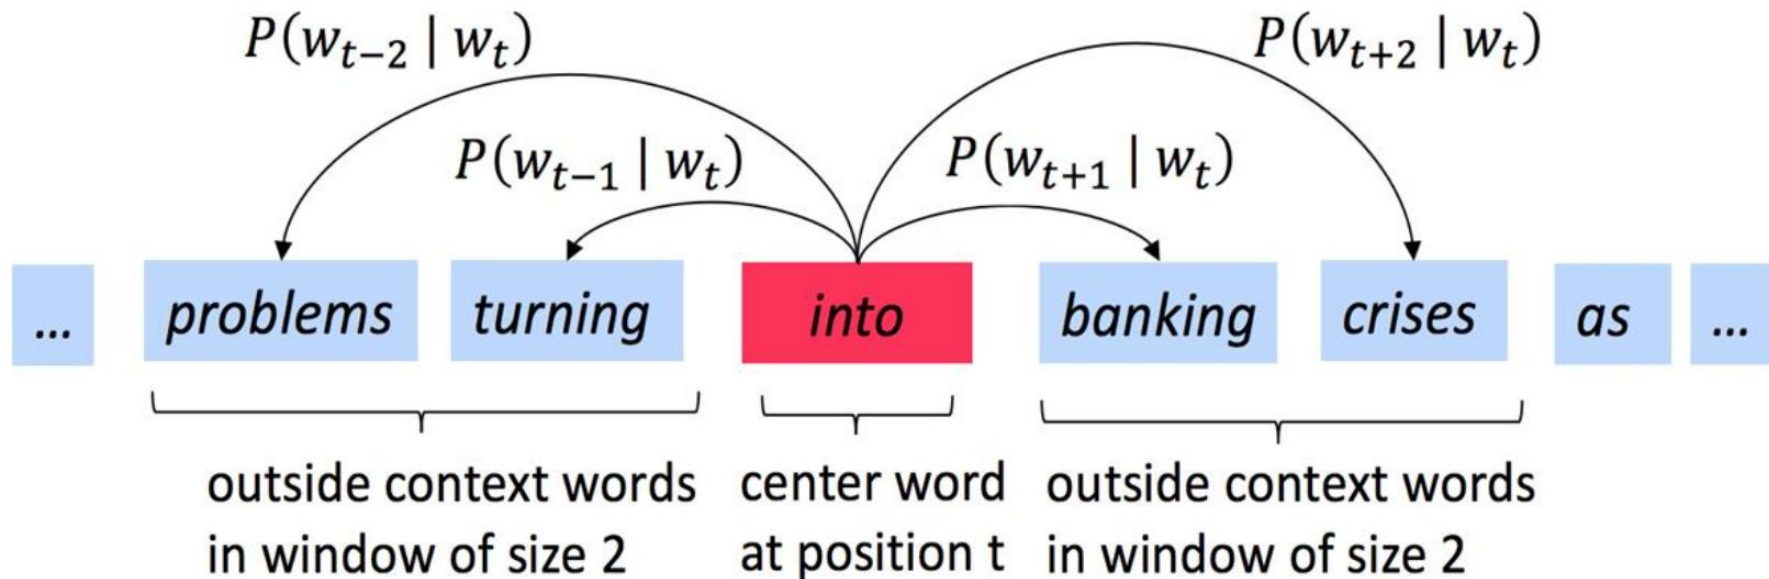


Source Text

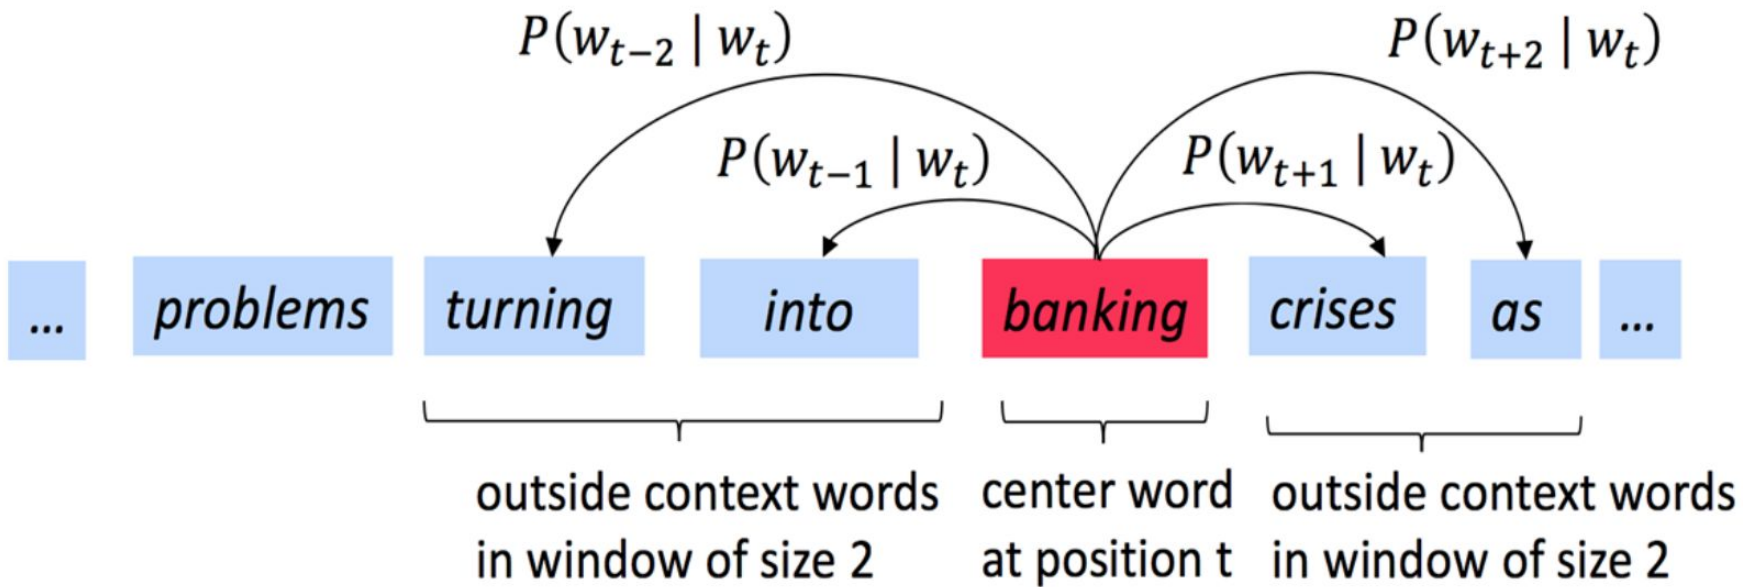
Training Samples

The quick brown fox jumps over the lazy dog. →	(the, quick) (the, brown)
The quick brown fox jumps over the lazy dog. →	(quick, the) (quick, brown) (quick, fox)
The quick brown fox jumps over the lazy dog. →	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
The quick brown fox jumps over the lazy dog. →	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

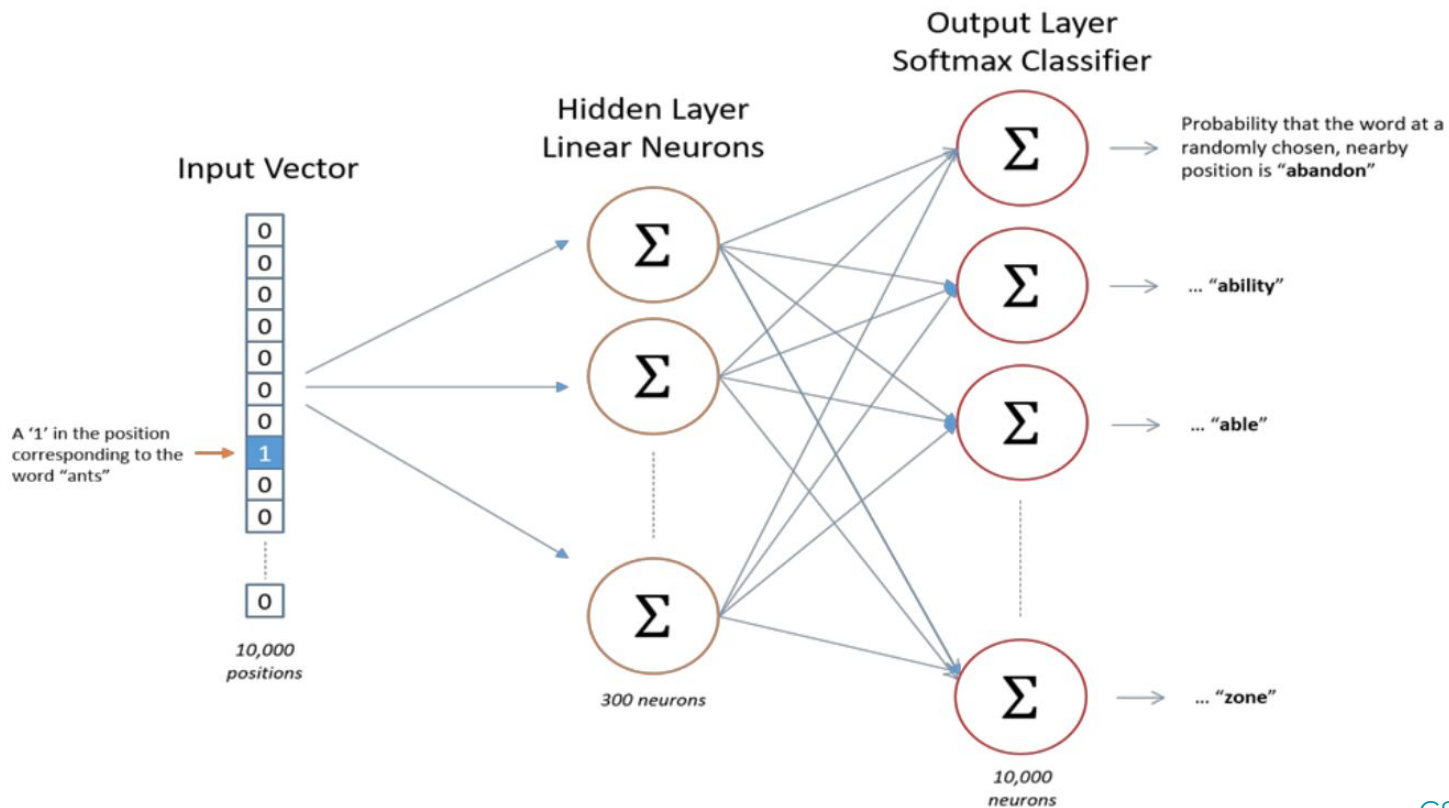
Word2vec



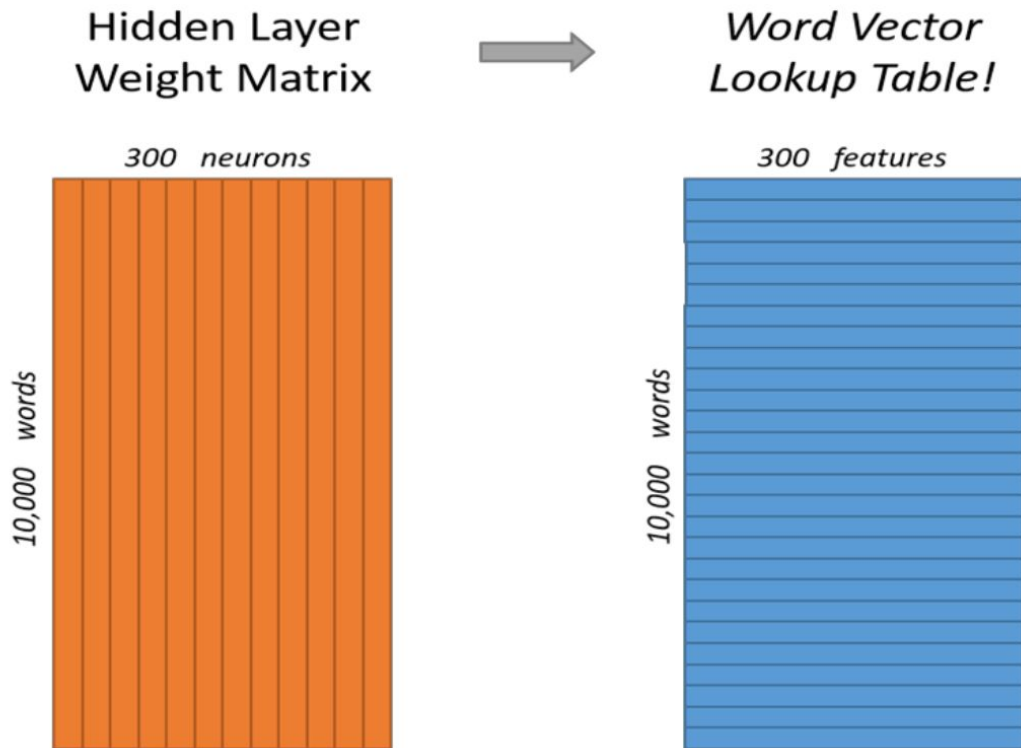
Word2vec



Word2vec



Word2vec





Word2vec

- Word vectors with 300 components
- Vocabulary of 10,000 words.
- Weight matrix with $300 \times 10,000 = 3$ million weights each!
- Training is too long and computationally expensive

HOW TO FIX IT?



Word2vec

Basic approaches:

1. Treating common word pairs or phrases as single “words” in their model.
2. Subsampling frequent words to decrease the number of training examples.
3. Modifying the optimization objective with a technique they called “Negative Sampling”, which causes each training sample to update only a small percentage of the model’s weights.




Word2vec: Subsampling

Subsampling frequent words.

w_i is the word, $z(w_i)$ is the fraction of this word in the text

Graph for $(\sqrt{x/0.001}+1)*0.001/x$



$P(w_i)$ is the probability of
keeping the word:

$$P(w_i) = \left(\sqrt{\frac{z(w_i)}{0.001}} + 1 \right) \cdot \frac{0.001}{z(w_i)}$$



Word2vec: Negative Sampling

- Error is computed only for a few words. All other words have zero error, so no updates by the backprop mechanism.
- More frequent words are selected to be negative samples more often. The probability for selecting a word is just its weight divided by the sum of weights for all words.

$$P(w_i) = \frac{f(w_i)^{3/4}}{\sum_{j=0}^n (f(w_j)^{3/4})}$$



Word2vec: two models

Continuous BOW (CBOW)

Predict center word from (bag of) context words

$$p(w_i | w_{i-h}, \dots, w_{i+h})$$

- Predicting one word each time
- Relatively fast


Skip-gram

Predict context ("outside") words (position independent) given center word

$$p(w_{i-h}, \dots, w_{i+h} | w_i)$$


- Predicting context by one word
- Much slower
- Better with infrequent words

Word2vec: Skip-gram





Word2vec: CBOW






Word2vec: Contrastive

Simplify the approaches above:

Consider only one **positive pair** of words and **a few negative pairs**.

$$\log \sigma(v_{w_o}^\top v_{w_t}) + \sum_{\substack{i=0 \\ w_i \in P_n(\omega)}}^k \log \sigma(-v_{w_i}^\top v_{w_t})$$

GloVe: visualizations






Word2vec: word analogies

King - man + woman = queen
↓ ↓ ↓ ↓
 x y y' $target$

$\cos(x - y + y', target) \rightarrow \max_{target}$



Summary



- Word vectors are simply vectors of numbers that represent the meaning of a word
- Approaches:
 - One-hot encoding
 - Bag-of-words models
 - Counts of word / context co-occurrences
 - TF-IDF
 - Predictions of context given word (skip-gram neural network models, e.g. word2vec)

Revise



1. NLP introduction
2. Text preprocessing
3. Feature extraction:
 - a. Bag-of-Words
 - b. Bag-of-Ngrammes
 - c. TF-IDF
4. Word embeddings
5. Word2vec

Thanks for attention!

Questions?

